

Data Import; Exploratory Data Analysis (EDA) in R

Based on lecture handouts originally written by Dr. David Gerard.

Learning Objectives

- Import data from CSV's,
- Working Directories
- Strategies for EDA
- [Data Import Cheat Sheet](#).
- [Readr Overview](#).

Part 1 (Data Import)

Working Directories

- The working directory is where R will look for and save things by default.
- When you specify to save a figure, save a file, or load some data, it will be with respect to the working directory.
- You can see where the current working directory is by `getwd()`, or by looking at the top of the console in RStudio.
- You can change the working directory by Session > Set Working Directory > Choose Directory. Or by CONTROL + SHIFT + H. Or you can use the `setwd()` command.
- A shortcut is to set the working directory to your source file location with Session > Set Working Directory > To Source File Location.
- When you read and write files/figures, you can then specify the path from the position of the working directory.
- Suppose we want to save the following figure:

```
suppressPackageStartupMessages(library(tidyverse))
```

```
## Warning: package 'tidyverse' was built under R version 4.0.5
```

```
## Warning: package 'tibble' was built under R version 4.0.5
```

```
## Warning: package 'tidyr' was built under R version 4.0.5
```

```
## Warning: package 'dplyr' was built under R version 4.0.4
```

```
## Warning: package 'forcats' was built under R version 4.0.5
```

```
data("mpg")  
pl <- ggplot(mpg, aes(x = hwy, y = cty)) +  
  geom_point()
```

- To save `p1` in the current folder, we would use:

```
ggsave(filename = "./my_saved_plot.pdf", plot = p1)
```

- The “.” means “the current folder”.
- To save `p1` in the folder one level up we would use:

```
ggsave(filename = "../my_saved_plot.pdf", plot = p1)
```

- The “..” means “go one level up”.
- If we are in the analysis folder, and we want to save `p1` in the output folder, we would use:

```
ggsave(filename = "../output/my_saved_plot.pdf", plot = p1)
```

- If we have a subfolder called “fig” within our current folder. We could save `p1` in “fig” with

```
ggsave(filename = "./fig/my_saved_plot.pdf", plot = p1)
```

- **NEVER USE ABSOLUTE PATHS.** For example, you should never start the path from “C” if you use Windows. This makes your code non-transferable to other users.

readr

- To read a CSV (comma-separated values) file into R, use the `read_csv()` function from the `readr` package.

```
suppressPackageStartupMessages(library(tidyverse))
heights <- read_csv(file = "./heights.csv")
```

```
##
## -- Column specification -----
## cols(
##   earn = col_double(),
##   height = col_double(),
##   sex = col_character(),
##   ed = col_double(),
##   age = col_double(),
##   race = col_character()
## )
```

- Use `read_tsv()` if columns are separated by tabs.
- Use `read_csv2()` if columns are separated by semicolons.
- Other file formats are listed in [RDS](#).
- First export the Excel spreadsheet as a CSV. Then read the CSV file into R.
- You are using colors to represent meaningful information in Excel? Don't.
 - Edit the data so that the information is encoded by a new variable.
- If you don't know the format ahead of time, use `read_lines()` to print the first few lines.

```
read_lines(file = "./heights.csv", n_max = 10)
```

```
## [1] "\"earn\", \"height\", \"sex\", \"ed\", \"age\", \"race\""
## [2] "50000,74.4244387818035, \"male\", 16, 45, \"white\""
## [3] "60000,65.5375428255647, \"female\", 16, 58, \"white\""
```

```
## [4] "30000,63.6291977374349,\"female\",16,29,\"white\""
## [5] "50000,63.1085616752971,\"female\",16,91,\"other\""
## [6] "51000,63.4024835710879,\"female\",17,39,\"white\""
## [7] "9000,64.3995075440034,\"female\",15,26,\"white\""
## [8] "29000,61.6563258264214,\"female\",12,49,\"white\""
## [9] "32000,72.6985437364783,\"male\",17,46,\"white\""
## [10] "2000,72.0394668497611,\"male\",15,21,\"hispanic\""
```

Special Considerations

- Always check your data immediately after importing it.
 - Check that the types are correct for each of the variables.
 - Check that the missing data were coded correctly.
 - Later on, when you notice something weird, consider that this might have resulted because of a problem during data import.

```
hate_crimes <- read_csv(file = "./hate_crimes2.csv")

##
## -- Column specification -----
## cols(
##   state = col_character(),
##   median_house_inc = col_double(),
##   share_unemp_seas = col_double(),
##   share_pop_metro = col_double(),
##   share_pop_hs = col_double(),
##   share_non_citizen = col_double(),
##   share_white_poverty = col_double(),
##   gini_index = col_double(),
##   share_non_white = col_double(),
##   share_vote_trump = col_double(),
##   hate_crimes_per_100k_splc = col_double(),
##   avg_hatecrimes_per_100k_fbi = col_double()
## )

summarize_all(hate_crimes, class)

## # A tibble: 1 x 12
##   state      median_house_inc share_unemp_seas share_pop_metro share_pop_hs
##   <chr>      <chr>              <chr>              <chr>              <chr>
## 1 character numeric          numeric          numeric          numeric
## # ... with 7 more variables: share_non_citizen <chr>,
## #   share_white_poverty <chr>, gini_index <chr>, share_non_white <chr>,
## #   share_vote_trump <chr>, hate_crimes_per_100k_splc <chr>,
## #   avg_hatecrimes_per_100k_fbi <chr>

summarize_all(hate_crimes, funs(sum(is.na(.)))) #sum all the NA's under each variable

## Warning: `funs()` was deprecated in dplyr 0.8.0.
## Please use a list of either functions or lambdas:
##
##   # Simple named list:
##   list(mean = mean, median = median)
##
```

```
## # Auto named with `tibble::lst()`:
## tibble::lst(mean, median)
##
## # Using lambdas
## list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))

## # A tibble: 1 x 12
## state median_house_inc share_unemp_seas share_pop_metro share_pop_hs
## <int> <int> <int> <int> <int>
## 1 0 0 0 0 0
## # ... with 7 more variables: share_non_citizen <int>,
## # share_white_poverty <int>, gini_index <int>, share_non_white <int>,
## # share_vote_trump <int>, hate_crimes_per_100k_splc <int>,
## # avg_hatecrimes_per_100k_fbi <int>
```

```
head(hate_crimes)
```

```
## # A tibble: 6 x 12
## state median_house_inc share_unemp_seas share_pop_metro share_pop_hs
## <chr> <dbl> <dbl> <dbl> <dbl>
## 1 Alabama 42278 0.06 0.64 0.821
## 2 Alaska 67629 0.064 0.63 0.914
## 3 Arizona 49254 0.063 0.9 0.842
## 4 Arkansas 44922 0.052 0.69 0.824
## 5 California 60487 0.059 0.97 0.806
## 6 Colorado 60940 0.04 0.8 0.893
## # ... with 7 more variables: share_non_citizen <dbl>,
## # share_white_poverty <dbl>, gini_index <dbl>, share_non_white <dbl>,
## # share_vote_trump <dbl>, hate_crimes_per_100k_splc <dbl>,
## # avg_hatecrimes_per_100k_fbi <dbl>
```

- Sometimes the files code missing data other than NA. For example, it's common to use periods ., or in some genomic settings they use -9 as missing.
- R won't know how to handle this without you telling it, so you'll have to know what the missing data encoding is and specify it with the `na` argument in `read_csv()`.
- readr will try to guess the type for each column (double, integer, character, logic, etc). Sometimes it guesses wrong. If it seems to be guessing wrong, use the `col_types` to explicitly specify the column types.
- Sometimes there are comments at the start of a data file. You can skip the first few lines before starting to read data with the `skip` argument.
- If the comments begin with a special character, you can use the `comment` argument.

Data Export

- You can write comma-separated and tab-separated files using `write_csv()`, `write_csv2()`, and `write_tsv()`.
- The defaults are usually fine.

Reading/Writing R Objects

- You can save and reload arbitrary R objects (data frames, matrices, lists, vectors) using `readRDS()` and `saveRDS()`.

Part 2 (Exploratory Data Analysis (EDA) in R)

We will use `ggplot2` which is a R package dedicated to data visualization.

General Strategies


- Plot the distribution of every variable.
- Plot the bivariate distribution of every pair of variables (to find which variables are associated).
- Color code by variables to try and see if relationships can be explained.
- Calculate lots of summary statistics.
- Look at missingness.
- Look at outliers.
- EDA is about **curiosity**. Ask *many* questions, use *many* plots, investigate *many* aspects of your data. This will let you hone in on the few *interesting* questions you want to pursue deeper.

```
library(tidyverse)
data("diamonds")
```

Distribution of Every Variable:

- Quantitative: Use a histogram.
 - Look for modality. Indicates multiple groups of units. What can explain the modes? Can any of the other variables explain the modes?
 - Are certain values more likely than other values?
 - Look for skew.
 - `geom_histogram()`
 - Mean, median, standard deviation, five number summary.

```
ggplot(data = diamonds, mapping = aes(x = carat)) +
  geom_histogram(bins = 500)
```



```
fivenum(diamonds$carat)
```

```
## [1] 0.20 0.40 0.70 1.04 5.01
```

```
mean(diamonds$carat)
```


```
## [1] 0.7979397
```

```
sd(diamonds$carat)
```

```
## [1] 0.4740112
```

- Categorical: Use a bar chart. Or just a table of *proportions* (`table()` then `prop.table()`).
 - Absolute counts are sometimes interesting, but usually you want to look at the proportion of observations in each category.
 - Is there a natural ordering of the categories (bad, medium, good)?
 - Why are some categories more represented than others?
 - `geom_bar()`, `geom_col()`
 - Proportion of observations within each group.

```
ggplot(diamonds, aes(x = color, y = )) +  
  geom_bar(aes(y = count / sum(..count..))) +  
  ylab("Proportion")
```



```
table(diamonds$color)
```

```
##
##      D      E      F      G      H      I      J
## 6775  9797  9542 11292  8304  5422  2808
```


```
prop.table(table(diamonds$color))
```

```
##
##           D           E           F           G           H           I           J
## 0.12560252 0.18162773 0.17690026 0.20934372 0.15394883 0.10051910 0.05205784
```

Bivariate Distribution of Every Pair of Variables

- Quantitative vs Quantitative: Use a scatterplot
 - Is the relationship linear? Quadratic? Exponential?
 - Logging is useful tool to make some associations linear. If the relationship is (i) monotonic and (ii) curved, then try logging the x-variable *if the x-variable is all positive*. If it is also (iii) more variable at larger y-values, then try logging the y-variable *instead of the x-variable if the y-variable is all positive*. Try logging both if you still see curvature *if both variables are all positive*.
 - Ask if an observed association can be explained by another variable?
 - Correlation coefficient (only appropriate if association is linear).
 - Kendall's tau (always appropriate).

```
ggplot(diamonds, aes(x = carat, y = price)) +
  geom_point() +
  scale_y_log10() +
  scale_x_log10()
```




```
cor(diamonds$carat, diamonds$price)
```

```
## [1] 0.9215913
```

```
## cor(diamonds$carat, diamonds$price, method = "kendall")
```

- Categorical vs Quantitative: Use a boxplot
 - For which levels of the categorical variable is the quantitative variable higher or lower?
 - For which levels is the quantitative variable more spread out?
 - Aggregated means, medians, standard deviations, quantiles

```
ggplot(diamonds, aes(x = color, y = price)) +  
  geom_boxplot() +  
  scale_y_log10()
```



```
diamonds %>%  
  mutate(logprice = log(price)) %>%  
  group_by(color) %>%
```



```
summarize(mean = mean(logprice),
          sd   = sd(logprice),
          median = median(logprice),
          Q1    = quantile(logprice, 0.25),
          Q3    = quantile(logprice, 0.75))
```

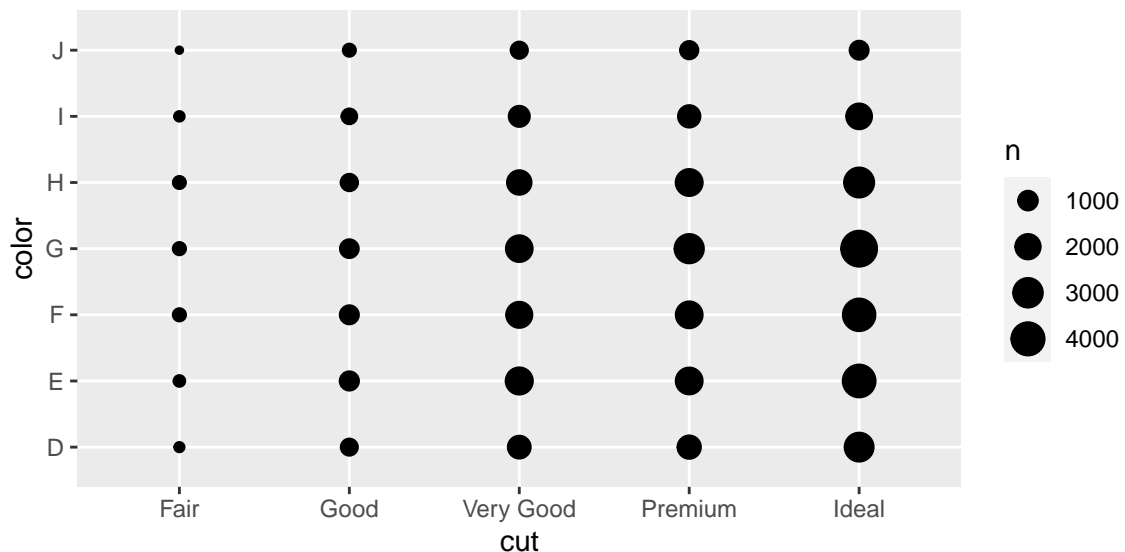
```
## # A tibble: 7 x 6
```

```
##   color mean    sd median    Q1    Q3
##   <ord> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 D      7.62 0.926   7.52  6.81  8.35
## 2 E      7.58 0.925   7.46  6.78  8.29
## 3 F      7.76 0.968   7.76  6.89  8.49
## 4 G      7.79 1.03    7.72  6.84  8.71
## 5 H      7.92 1.06    8.15  6.89  8.70
## 6 I      8.02 1.11    8.22  7.02  8.88
## 7 J      8.15 1.04    8.35  7.53  8.95
```

- Categorical vs Categorical: Use a mosaic plot or a count plot
 - For which pairs of values of the categorical variables are there the most number of units?


```
## Only gives you the bivariate distribution
```

```
ggplot(diamonds, aes(x = cut, y = color)) +
  geom_count()
```



```
## Gives you the conditional distributions of color given cut
```

```
ggplot(diamonds, aes(x = cut, fill = color)) +
  geom_bar(position = "fill")
```



```
## Gives you the conditional distributions of cut given color
ggplot(diamonds, aes(x = color, fill = cut)) +
  geom_bar(position = "fill")
```

